

Если прочитать описание данных, можно увидеть следующее:

shift - на сколько недель значения признаков отстают от недели, записанной в week (отставание признаков нужно чтобы научиться прогнозировать продажи на 2 или 3 недели вперед)

Действительно напрашивается какое-то преобразование датасета, так как в данный момент, если посмотреть на датасет train, видно, что у нас одинаковые целевые переменные, например 110776, разные shift, и абсолютно разный набор f-ок.

Довольно странно ожидать хорошего качества регрессии, если у нас при абсолютно разных наборах X одинаковые значения y.

Можно сдвинуть признаки к соответствующим неделям, чтобы не было запоздания, но это получится сделать не у всех, их мы в данный момент просто выкинем. В тестовом датасете получится сделать такой сдвиг у 2/3 записей, я обучал отдельно модель1 для записей со сдвигом (2/3) (конечно же на преобразованном датасете) и отдельно модель2 для предсказания записей без сдвига (1/3). Я не очень заморачивался с подбором параметром, возможно, при хорошем переборе, можно добиться улучшения на 0.5-1.0 по метрике smape. Мой финальный результат 9.7130.

```
In [1]: import numpy as np
import pandas as pd
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error
from sklearn.linear_model import LinearRegression, Lasso
from collections import Counter
import xgboost
import matplotlib.pyplot as plt
from sklearn import cross_validation, grid_search
import seaborn as sns
from tqdm import tqdm
from sklearn.utils import shuffle
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

```
/Users/daniilkorbut/anaconda/lib/python3.5/site-packages/sklearn/cross_validation.py:44: DeprecationWarning:
  This module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored
  classes and functions are moved. Also note that the interface of the new CV iterators are different from
  that of this module. This module will be removed in 0.20.
  "This module will be removed in 0.20.", DeprecationWarning)
/Users/daniilkorbut/anaconda/lib/python3.5/site-packages/sklearn/grid_search.py:43: DeprecationWarning: This
  module was deprecated in version 0.18 in favor of the model_selection module into which all the refactored c
  lasses and functions are moved. This module will be removed in 0.20.
  DeprecationWarning)
```

```
In [2]: train = pd.read_csv("train.tsv")
test = pd.read_csv("test.tsv")
sample_submission = pd.read_csv("sample_submission.tsv")
```

```
In [3]: test.head()
```

Out[3]:

	Num	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f51	f52	f53	f54	f55
0	348622	2015	3	3	20447918	960.0	820.0	1128.0	1801.0	1045.0	...	1510.0	580.0	969.0	1635.0	895.0
1	348623	2015	3	3	20447902	9086.0	12585.0	11595.0	9685.0	12917.0	...	22055.0	14235.0	21195.0	18280.0	182.0
2	348624	2015	3	3	20447732	115087.0	147287.0	176065.0	143105.0	202069.0	...	302165.0	162232.0	221622.0	256605.0	240.0
3	348625	2015	3	3	20443951	20900.0	24420.0	27068.0	20460.0	25580.0	...	39055.0	14445.0	22450.0	22093.0	311.0
4	348626	2015	3	3	20443944	4430.0	5864.0	3310.0	1853.0	2836.0	...	120.0	130.0	60.0	30.0	50.0

5 rows × 65 columns

```
In [4]: train.head()
```

Out[4]:

	Num	y	year	week	shift	item_id	f1	f2	f3	f4	...	f51	f52	f53	f54	f55	f56
0	0	123438	2012	52	1	20442076	4915.0	38056.0	40185.0	45733.0	...	39423.0	41765.0	52590.0	31452.0	44420.0	41865.0
1	1	58410	2012	52	1	20441997	2230.0	18817.0	20110.0	26368.0	...	22830.0	25230.0	27850.0	21390.0	27090.0	23170.0
2	2	163930	2012	52	1	20441990	5695.0	47480.0	47619.0	89708.0	...	14930.0	44290.0	46412.0	29320.0	21140.0	28406.0
3	3	53902	2012	52	1	20441989	1995.0	17146.0	20066.0	27070.0	...	15120.0	12480.0	19780.0	7990.0	8230.0	10650.0
4	4	105970	2012	52	1	20441988	6515.0	49262.0	50045.0	95167.0	...	18872.0	19328.0	37168.0	13570.0	19760.0	20208.0


5 rows × 66 columns

```
In [5]: def smape(y_pred, y_true):
        return (100 / y_pred.shape[0]) * np.sum(np.abs(y_pred - y_true) / ((np.abs(y_pred) + np.abs(y_true)) / 2))
```

Давайте посмотрим на забавную матрицу корреляций признаков

```
In [6]: corrmatrix = train.corr()
f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corrmatrix, vmax=.8, square=True)
```

```
Out[6]: <matplotlib.axes._subplots.AxesSubplot at 0x115dcc5f8>
```



Также стоит зметить, что признаки f1 и f31 полностью совпадают, f2 и f32 и тд., нет смысла включать их в модель

```
In [7]: print (np.array(train["f1"]))
print (np.array(train["f31"]))
print (np.array(train["f2"]))
print (np.array(train["f32"]))
res = 0
for i in range(1, 31):
    if (np.sum(np.array(train["f"+str(i)])) == np.array(train["f"+str(i+30)])) == train.shape[0]):
        res += 1
print(res)
```

```
[ 4915.  2230.  5695. ...,  515.  5597.    10.]
[ 4915.  2230.  5695. ...,  515.  5597.    10.]
[ 38056.  18817.  47480. ...,   795.   5595.     0.]
[ 38056.  18817.  47480. ...,   795.   5595.     0.]
30
```

```
In [8]: for i in range(31, 61):
        del train["f" + str(i)]
        del test["f" + str(i)]
```

```
In [9]: print (train.shape)
train.head()
```

```
(72457, 36)
```

```
Out[9]:
```

	Num	y	year	week	shift	item_id	f1	f2	f3	f4	...	f21	f22	f23	f24	f25	f26
0	0	123438	2012	52	1	20442076	4915.0	38056.0	40185.0	45733.0	...	39423.0	41765.0	52590.0	31452.0	44420.0	41865.0
1	1	58410	2012	52	1	20441997	2230.0	18817.0	20110.0	26368.0	...	22830.0	25230.0	27850.0	21390.0	27090.0	23170.0
2	2	163930	2012	52	1	20441990	5695.0	47480.0	47619.0	89708.0	...	14930.0	44290.0	46412.0	29320.0	21140.0	28406.0
3	3	53902	2012	52	1	20441989	1995.0	17146.0	20066.0	27070.0	...	15120.0	12480.0	19780.0	7990.0	8230.0	10650.0
4	4	105970	2012	52	1	20441988	6515.0	49262.0	50045.0	95167.0	...	18872.0	19328.0	37168.0	13570.0	19760.0	20208.0

```
5 rows × 36 columns
```

```
In [10]: print (test.shape)
         test.head()
```

```
(2016, 35)
```

```
Out[10]:
```

	Num	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24	f25
0	348622	2015	3	3	20447918	960.0	820.0	1128.0	1801.0	1045.0	...	1510.0	580.0	969.0	1635.0	895
1	348623	2015	3	3	20447902	9086.0	12585.0	11595.0	9685.0	12917.0	...	22055.0	14235.0	21195.0	18280.0	182
2	348624	2015	3	3	20447732	115087.0	147287.0	176065.0	143105.0	202069.0	...	302165.0	162232.0	221622.0	256605.0	240
3	348625	2015	3	3	20443951	20900.0	24420.0	27068.0	20460.0	25580.0	...	39055.0	14445.0	22450.0	22093.0	311
4	348626	2015	3	3	20443944	4430.0	5864.0	3310.0	1853.0	2836.0	...	120.0	130.0	60.0	30.0	50.0

```
5 rows × 35 columns
```

Я уже знаю, что ждёт нас дальше, так что давайте просто обучим RandomForest на неизменённом датасете (именно он показал самое хорошее качество из первых 3 посылок)

```
In [11]: model_1 = RandomForestRegressor(n_estimators=50, min_samples_leaf=2, min_samples_split=4)
```

```
In [12]: train1 = train.head(50000)
         train2 = train.tail(train.shape[0]-50000)
```

```
In [13]: X1 = train1.drop(['Num', 'y'], axis=1)
         y1 = train1['y']
         X2 = train2.drop(['Num', 'y'], axis=1)
         y2 = train2['y']
         y1 = np.array(y1)
         y2 = np.array(y2)
```

```
In [14]: %%time  
         model_1.fit(X1, y1)
```

```
CPU times: user 1min 8s, sys: 954 ms, total: 1min 9s  
Wall time: 1min 15s
```

```
Out[14]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,  
                               max_features='auto', max_leaf_nodes=None,  
                               min_impurity_split=1e-07, min_samples_leaf=2,  
                               min_samples_split=4, min_weight_fraction_leaf=0.0,  
                               n_estimators=50, n_jobs=1, oob_score=False, random_state=None,  
                               verbose=0, warm_start=False)
```

```
In [15]: res = model_1.predict(X2)  
         smape(res, y2)
```

```
Out[15]: 28.648295586100343
```

```
In [16]: train = train.sort(['item_id', 'year', 'week', 'shift'])  
         print (train.shape)  
  
         (72457, 36)
```

```
In [17]: del train['Num']
```

```
In [18]: converted_train = train
```

Посмотрим, сколько недель в 2013 и 2014 годах, это нам поможет при изменении датасета

```
In [19]: converted_train[converted_train["year"] == 2013].tail(5)
```

```
Out[19]:
```

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
34531	223657	2013	51	2	20453563	229136.0	237506.0	295042.0	210294.0	132000.0	...	270833.0	101837.0	116481.0	117941.0
34086	223657	2013	51	3	20453563	180591.0	229136.0	237506.0	295042.0	210294.0	...	167200.0	270833.0	101837.0	116481.0
35655	329673	2013	52	1	20453563	295042.0	210294.0	132000.0	237890.0	218525.0	...	116481.0	117941.0	212682.0	54410.0
35205	329673	2013	52	2	20453563	237506.0	295042.0	210294.0	132000.0	237890.0	...	101837.0	116481.0	117941.0	212682.0
34756	329673	2013	52	3	20453563	229136.0	237506.0	295042.0	210294.0	132000.0	...	270833.0	101837.0	116481.0	117941.0

5 rows × 35 columns

```
In [20]: converted_train[converted_train["year"] == 2014].tail(5)
```

```
Out[20]:
```

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
70077	171081	2014	52	2	20453563	125906.0	151188.0	120173.0	103150.0	128910.0	...	68340.0	66660.0	76850.0	107219.0
69625	171081	2014	52	3	20453563	112293.0	125906.0	151188.0	120173.0	103150.0	...	155135.0	68340.0	66660.0	76850.0
71197	168919	2014	53	1	20453563	120173.0	103150.0	128910.0	119210.0	118701.0	...	76850.0	107219.0	55220.0	83429.0
70751	168919	2014	53	2	20453563	151188.0	120173.0	103150.0	128910.0	119210.0	...	66660.0	76850.0	107219.0	55220.0
70302	168919	2014	53	3	20453563	125906.0	151188.0	120173.0	103150.0	128910.0	...	68340.0	66660.0	76850.0	107219.0

5 rows × 35 columns


```
In [21]: converted_train_m = converted_train.as_matrix()
print (converted_train_m.shape)
converted_train_m
```

```
(72457, 35)
```

```
Out[21]: array([[ 3.40138000e+05,  2.01200000e+03,  5.20000000e+01, ...,
                8.87890000e+04,  1.08817000e+05,  1.10228000e+05],
               [ 1.10776000e+05,  2.01300000e+03,  1.00000000e+00, ...,
                1.08817000e+05,  1.10228000e+05,  2.11266000e+05],
               [ 1.10776000e+05,  2.01300000e+03,  1.00000000e+00, ...,
                8.87890000e+04,  1.08817000e+05,  1.10228000e+05],
               ...,
               [ 1.26618000e+05,  2.01500000e+03,  2.00000000e+00, ...,
                1.06262000e+05,  1.04919000e+05,  1.83820000e+04],
               [ 1.26618000e+05,  2.01500000e+03,  2.00000000e+00, ...,
                8.19260000e+04,  1.06262000e+05,  1.04919000e+05],
               [ 1.26618000e+05,  2.01500000e+03,  2.00000000e+00, ...,
                7.41740000e+04,  8.19260000e+04,  1.06262000e+05]])
```

```
In [22]: modified = np.zeros(converted_train_m.shape[0])
```

```
In [23]: train.head()
```

```
Out[23]:
```

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
150	340138	2012	52	1	20427028	9700.0	92038.0	111363.0	146522.0	174408.0	...	93799.0	113945.0	156656.0	69913.0
841	110776	2013	1	1	20427028	92038.0	111363.0	146522.0	174408.0	111337.0	...	113945.0	156656.0	69913.0	96291.0
380	110776	2013	1	2	20427028	9700.0	92038.0	111363.0	146522.0	174408.0	...	93799.0	113945.0	156656.0	69913.0
1536	70915	2013	2	1	20427028	111363.0	146522.0	174408.0	111337.0	124149.0	...	156656.0	69913.0	96291.0	112953.0
1072	70915	2013	2	2	20427028	92038.0	111363.0	146522.0	174408.0	111337.0	...	113945.0	156656.0	69913.0	96291.0

```
5 rows × 35 columns
```

```

In [24]: year_pos = 1
         week_pos = 2
         shift_pos = 3
         item_pos = 4
         for i in tqdm(range(converted_train_m.shape[0])):
             year1 = converted_train_m[i][year_pos]
             week1 = converted_train_m[i][week_pos]
             shift1 = converted_train_m[i][shift_pos]
             item1 = converted_train_m[i][item_pos]
             week2 = week1 - shift1
             year2 = year1
             #if (i == 6):
             #    print (week1, year1, item1, shift1)
             if (week2 <= 0):
                 if (year1 == 2014):
                     week2 = 52 - week1
                 if (year2 == 2015):
                     week2 = 53 - week1
                 year2 = year1 - 1
             for j in range(i, 0, -1):
                 #if (i == 6):
                 #    print (j, ":", converted_train_m[j][year_pos], converted_train_m[j][week_pos], converted_train_m[j][item_pos], converted_train_m[j][shift_pos])
                 if (converted_train_m[j][year_pos] == year2 and converted_train_m[j][week_pos] == week2 and
                     converted_train_m[j][item_pos] == item1 and converted_train_m[j][shift_pos] == shift1):
                     if (modified[j] == 0):
                         converted_train_m[j][5:35] = converted_train_m[i][5:35]
                         modified[j] = 1
                     break

```

100%|██████████| 72457/72457 [01:11<00:00, 1013.90it/s]

```

In [25]: print (len(modified))
         print (np.sum(modified))

```

72457
68763.0

```

In [26]: corrected_train_m = converted_train_m

```

```
In [27]: cols = train.columns
cols
```

```
Out[27]: Index(['y', 'year', 'week', 'shift', 'item_id', 'f1', 'f2', 'f3', 'f4', 'f5',
               'f6', 'f7', 'f8', 'f9', 'f10', 'f11', 'f12', 'f13', 'f14', 'f15', 'f16',
               'f17', 'f18', 'f19', 'f20', 'f21', 'f22', 'f23', 'f24', 'f25', 'f26',
               'f27', 'f28', 'f29', 'f30'],
              dtype='object')
```

```
In [28]: new_dataset_train = pd.DataFrame(converted_train_m, columns=cols)
```

```
In [29]: new_dataset_train.head()
```

```
Out[29]:
```

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
0	340138.0	2012.0	52.0	1.0	20427028.0	9700.0	92038.0	111363.0	146522.0	174408.0	...	93799.0	113945.0	156656.0	69913.0
1	110776.0	2013.0	1.0	1.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	156656.0	69913.0	96291.0	112953.0
2	110776.0	2013.0	1.0	2.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	156656.0	69913.0	96291.0	112953.0
3	70915.0	2013.0	2.0	1.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	69913.0	96291.0	112953.0	164259.0
4	70915.0	2013.0	2.0	2.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	69913.0	96291.0	112953.0	164259.0

5 rows × 35 columns

```
In [30]: new_dataset_train["modi"] = modified
new_dataset_train.head()
```

Out[30]:

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f22	f23	f24	f25
0	340138.0	2012.0	52.0	1.0	20427028.0	9700.0	92038.0	111363.0	146522.0	174408.0	...	113945.0	156656.0	69913.0	96291.0
1	110776.0	2013.0	1.0	1.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	69913.0	96291.0	112953.0	164259.0
2	110776.0	2013.0	1.0	2.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	69913.0	96291.0	112953.0	164259.0
3	70915.0	2013.0	2.0	1.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	96291.0	112953.0	164259.0	88789.0
4	70915.0	2013.0	2.0	2.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	96291.0	112953.0	164259.0	88789.0

5 rows × 36 columns

```
In [31]: new_dataset_train = new_dataset_train[new_dataset_train["modi"] == 1]
del new_dataset_train['modi']
new_dataset_train.head()
```

Out[31]:

	y	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
1	110776.0	2013.0	1.0	1.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	156656.0	69913.0	96291.0	112953.0
2	110776.0	2013.0	1.0	2.0	20427028.0	111363.0	146522.0	174408.0	111337.0	124149.0	...	156656.0	69913.0	96291.0	112953.0
3	70915.0	2013.0	2.0	1.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	69913.0	96291.0	112953.0	164259.0
4	70915.0	2013.0	2.0	2.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	69913.0	96291.0	112953.0	164259.0
5	70915.0	2013.0	2.0	3.0	20427028.0	146522.0	174408.0	111337.0	124149.0	119286.0	...	69913.0	96291.0	112953.0	164259.0

5 rows × 35 columns

Теперь поработаем с тестовым датасетом

```
In [32]: print (test.shape)
test.sort(['item_id', 'year', 'week', 'shift']).head()
```

```
(2016, 35)
```

```
Out[32]:
```

	Num	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f21	f22	f23	f24
822	350335	2015	3	1	20427028	112386.0	103264.0	98742.0	108584.0	102320.0	...	95788.0	81425.0	105634.0	96496.0
370	349663	2015	3	2	20427028	107536.0	112386.0	103264.0	98742.0	108584.0	...	61980.0	95788.0	81425.0	105634.0
144	348766	2015	3	3	20427028	92012.0	107536.0	112386.0	103264.0	98742.0	...	126866.0	61980.0	95788.0	81425.0
1496	351009	2015	4	1	20427028	103264.0	98742.0	108584.0	102320.0	122395.0	...	81425.0	105634.0	96496.0	82307.0
1050	350563	2015	4	2	20427028	112386.0	103264.0	98742.0	108584.0	102320.0	...	95788.0	81425.0	105634.0	96496.0

```
5 rows × 35 columns
```

```
In [33]: test_m = test.as_matrix()
print (test_m.shape)
test_m
```

```
(2016, 35)
```

```
Out[33]: array([[ 3.48622000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  1.02000000e+03,  1.29300000e+03,  1.29000000e+03],
                [ 3.48623000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  1.83200000e+04,  2.41160000e+04,  2.13070000e+04],
                [ 3.48624000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  2.45652000e+05,  2.86179000e+05,  2.85904000e+05],
                ...,
                [ 3.51526000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  1.13270000e+04,  1.07320000e+04,  1.57090000e+04],
                [ 3.51527000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  3.37000000e+02,  1.95000000e+02,  2.00000000e+02],
                [ 3.51528000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  2.59700000e+03,  3.98100000e+03,  4.48900000e+03]])
```

```
In [34]: modified = np.zeros(test_m.shape[0])
```

```

In [35]: year_pos = 1
        week_pos = 2
        shift_pos = 3
        item_pos = 4
        for i in tqdm(range(test_m.shape[0])):
            year1 = test_m[i][year_pos]
            week1 = test_m[i][week_pos]
            shift1 = test_m[i][shift_pos]
            item1 = test_m[i][item_pos]
            week2 = week1 - shift1
            year2 = year1
            #if (i == 3):
            #    print (week1, year1, item1, shift1)
            if (week2 <= 0):
                if (year1 == 2014):
                    week2 = 52 - week1
                if (year2 == 2015):
                    week2 = 53 - week1
                year2 = year1 - 1
            for j in range(0, test_m.shape[0]):
                #if (i == 3):
                #    print (j, ":", test_m[j][year_pos], test_m[j][week_pos], test_m[j][item_pos])
                if (test_m[j][year_pos] == year2 and test_m[j][week_pos] == week2 and
                    test_m[j][item_pos] == item1):
                    if (modified[j] == 0):
                        test_m[j][5:65] = test_m[i][5:65]
                        modified[j] = 1

```

100%|██████████| 2016/2016 [00:05<00:00, 340.82it/s]

```
In [36]: test_m
```

```
Out[36]: array([[ 3.48622000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  4.22000000e+02,  1.69000000e+03,  9.30000000e+02],
                [ 3.48623000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  6.66100000e+03,  2.06100000e+04,  1.67910000e+04],
                [ 3.48624000e+05,  2.01500000e+03,  3.00000000e+00, ...,
                  5.48460000e+04,  1.71320000e+05,  1.84820000e+05],
                ...,
                [ 3.51526000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  1.13270000e+04,  1.07320000e+04,  1.57090000e+04],
                [ 3.51527000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  3.37000000e+02,  1.95000000e+02,  2.00000000e+02],
                [ 3.51528000e+05,  2.01500000e+03,  5.00000000e+00, ...,
                  2.59700000e+03,  3.98100000e+03,  4.48900000e+03]])
```

```
In [37]: print (len(modified))
print (np.sum(modified))
```

```
2016
1335.0
```

Видим, что около 2/3 мы смогли преобразовать, как мы делали это в train, но для последней недели нет следующей, из которой мы бы могли взять информацию

```
In [38]: cols = test.columns
cols
```

```
Out[38]: Index(['Num', 'year', 'week', 'shift', 'item_id', 'f1', 'f2', 'f3', 'f4', 'f5',
               'f6', 'f7', 'f8', 'f9', 'f10', 'f11', 'f12', 'f13', 'f14', 'f15', 'f16',
               'f17', 'f18', 'f19', 'f20', 'f21', 'f22', 'f23', 'f24', 'f25', 'f26',
               'f27', 'f28', 'f29', 'f30'],
              dtype='object')
```

```
In [39]: new_dataset_test = pd.DataFrame(test_m, columns=cols)
```

```
In [40]: new_dataset_test["modi"] = modified
```

```
In [41]: new_dataset_test.head()
```

```
Out[41]:
```

	Num	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f22	f23	f24	f25
0	348622.0	2015.0	3.0	3.0	20447918.0	1801.0	1045.0	1190.0	1310.0	1111.0	...	895.0	2140.0	1182.0	1020.0
1	348623.0	2015.0	3.0	3.0	20447902.0	9685.0	12917.0	11157.0	12105.0	11765.0	...	18270.0	15851.0	16920.0	18320.0
2	348624.0	2015.0	3.0	3.0	20447732.0	143105.0	202069.0	216920.0	236053.0	257198.0	...	240047.0	236630.0	206697.0	245652.0
3	348625.0	2015.0	3.0	3.0	20443951.0	20460.0	25580.0	27135.0	24059.0	27335.0	...	31175.0	23355.0	15358.0	18930.0
4	348626.0	2015.0	3.0	3.0	20443944.0	1853.0	2836.0	2807.0	2550.0	3951.0	...	50.0	20.0	20.0	30.0

5 rows × 36 columns

Мы уже обучили в самом начале ноутбука модель под неизменённый датасет, и теперь необходимо обучить модель на изменённый, в тестовой выборке будем предсказывать целевую переменную для неизменённых строк с помощью первой модели, а для изменённых с помощью второй

```
In [42]: model_2 = RandomForestRegressor(n_estimators=100, min_samples_leaf=2, min_samples_split=2)
```

```
In [43]: new_dataset_train.shape
```

```
Out[43]: (68763, 35)
```

```
In [44]: X1 = new_dataset_train.head(50000).drop(['y'], axis=1)
y1 = new_dataset_train.head(50000)['y']
X2 = new_dataset_train.tail(18000).drop(['y'], axis=1)
y2 = new_dataset_train.tail(18000)['y']
y1 = np.array(y1)
y2 = np.array(y2)
```



```
In [45]: model_2.fit(X1, y1)
```

```
Out[45]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,  
                                max_features='auto', max_leaf_nodes=None,  
                                min_impurity_split=1e-07, min_samples_leaf=2,  
                                min_samples_split=2, min_weight_fraction_leaf=0.0,  
                                n_estimators=100, n_jobs=1, oob_score=False, random_state=None,  
                                verbose=0, warm_start=False)
```

```
In [46]: res = model_2.predict(X2)  
         smape(res, y2)
```

```
Out[46]: 1.0072707708915898
```

```

In [54]: %%time
results = []
for n_e in [10, 20, 40]:
    for m_s_l in [1, 3, 5]:
        for m_s_s in [2, 3]:
            tmp = []
            for i in range(5):
                trainnnnn = shuffle(new_dataset_train)
                train1 = trainnnnn.head(45000)
                train2 = trainnnnn.tail(train.shape[0]-45000)
                X1 = train1.drop(['y'], axis=1)
                y1 = train1['y']
                X2 = train2.drop(['y'], axis=1)
                y2 = train2['y']
                y1 = np.array(y1)
                y2 = np.array(y2)
                model = RandomForestRegressor(n_estimators=n_e, min_samples_leaf=m_s_l, min_samples_split=m_s_s)
                model.fit(X1, y1)
                res = model.predict(X2)
                tmp.append(sape(res, y2))
            results.append((np.mean(tmp), n_e, m_s_l, m_s_s))

(.0)
129
130     def __call__(self):
--> 131         return [func(*args, **kwargs) for func, args, kwargs in self.items]
132
133     def __len__(self):

/Users/daniilkorbut/anaconda/lib/python3.5/site-packages/sklearn/ensemble/forest.py in
_parallel_build_trees(tree, forest, X, y, sample_weight, tree_idx, n_trees, verbose, class_weight)
118         curr_sample_weight *= compute_sample_weight('balanced', y, indices)
119
--> 120         tree.fit(X, y, sample_weight=curr_sample_weight, check_input=False)
121     else:
122         tree.fit(X, y, sample_weight=sample_weight, check_input=False)

/Users/daniilkorbut/anaconda/lib/python3.5/site-packages/sklearn/tree/tree.py in fit(self, X, y, sample_weight,
check_input, X_idx_sorted)
1027         sample_weight=sample_weight,
1028         check_input=check_input,
~ 1029         X_idx_sorted=X_idx_sorted)

```

```
In [48]: new_dataset_test.head()
```

```
Out[48]:
```

	Num	year	week	shift	item_id	f1	f2	f3	f4	f5	...	f22	f23	f24	f25
0	348622.0	2015.0	3.0	3.0	20447918.0	1801.0	1045.0	1190.0	1310.0	1111.0	...	895.0	2140.0	1182.0	1020.0
1	348623.0	2015.0	3.0	3.0	20447902.0	9685.0	12917.0	11157.0	12105.0	11765.0	...	18270.0	15851.0	16920.0	18320.0
2	348624.0	2015.0	3.0	3.0	20447732.0	143105.0	202069.0	216920.0	236053.0	257198.0	...	240047.0	236630.0	206697.0	245652.0
3	348625.0	2015.0	3.0	3.0	20443951.0	20460.0	25580.0	27135.0	24059.0	27335.0	...	31175.0	23355.0	15358.0	18930.0
4	348626.0	2015.0	3.0	3.0	20443944.0	1853.0	2836.0	2807.0	2550.0	3951.0	...	50.0	20.0	20.0	30.0

5 rows × 36 columns

```
In [49]: np.sum(new_dataset_test['modi'] == modified)
```

```
Out[49]: 2016
```

```
In [50]: del new_dataset_test['modi']
```

```
In [51]: del new_dataset_test['Num']
```

```
In [52]: new_dataset_test.head()
```

```
Out[52]:
```

	year	week	shift	item_id	f1	f2	f3	f4	f5	f6	...	f21	f22	f23	f24
0	2015.0	3.0	3.0	20447918.0	1801.0	1045.0	1190.0	1310.0	1111.0	870.0	...	1635.0	895.0	2140.0	1182.0
1	2015.0	3.0	3.0	20447902.0	9685.0	12917.0	11157.0	12105.0	11765.0	10823.0	...	18280.0	18270.0	15851.0	16920.0
2	2015.0	3.0	3.0	20447732.0	143105.0	202069.0	216920.0	236053.0	257198.0	267316.0	...	256605.0	240047.0	236630.0	206697.0
3	2015.0	3.0	3.0	20443951.0	20460.0	25580.0	27135.0	24059.0	27335.0	23271.0	...	22093.0	31175.0	23355.0	15358.0
4	2015.0	3.0	3.0	20443944.0	1853.0	2836.0	2807.0	2550.0	3951.0	3398.0	...	30.0	50.0	20.0	20.0

5 rows × 34 columns

```
In [53]: new_dataset_test.iloc[0,:]
```

```
Out[53]: year          2015.0  
         week           3.0  
         shift          3.0  
         item_id      20447918.0  
         f1            1801.0  
         f2            1045.0  
         f3            1190.0  
         f4            1310.0  
         f5            1111.0  
         f6             870.0  
         f7            1180.0  
         f8             960.0  
         f9            1301.0  
         f10           1367.0  
         f11           1960.0  
         f12           6460.0  
         f13           1272.0  
         f14           1030.0  
         f15           1092.0  
         f16           1390.0  
         f17           1710.0  
         f18           1510.0  
         f19            580.0  
         f20            969.0  
         f21           1635.0  
         f22            895.0  
         f23           2140.0  
         f24           1182.0  
         f25           1020.0  
         f26           1293.0  
         f27           1290.0  
         f28            422.0  
         f29           1690.0  
         f30            930.0  
         Name: 0, dtype: float64
```

```
In [ ]:
```

```
In [54]: %%time

preds = []
for i in range(len(modified)):
    if (modified[i]):
        res = model_2.predict(np.array(new_dataset_test.iloc[i,:]))
        preds.append(res)
    else:
        res = model_2.predict(np.array(new_dataset_test.iloc[i,:]))
        preds.append(res)

print (len(preds))
print (len(sample_submission))
```

```
2016
2016
CPU times: user 13.7 s, sys: 278 ms, total: 14 s
Wall time: 15 s
```

```
In [55]: preds = np.array(preds)
preds
```

```
Out[55]: array([[ 1497.492      ],
 [ 27029.30716667],
 [ 297065.0302619 ],
 ...,
 [ 25297.52266667],
 [    322.        ],
 [ 7222.58766667]])
```

```
In [56]: preds = np.ravel(preds)
preds
```

```
Out[56]: array([ 1497.492      , 27029.30716667, 297065.0302619 , ...,
 25297.52266667,    322.        , 7222.58766667])
```

```
In [57]: sample_submission['y'] = preds
```

```
In [58]: sample_submission.head(5)
```

```
Out[58]:
```

	Num	y
0	348622	1497.492000
1	348623	27029.307167
2	348624	297065.030262
3	348625	29172.000286
4	348626	15.993333

```
In [59]: print (sample_submission[sample_submission['y'] < 0])
```

```
Empty DataFrame  
Columns: [Num, y]  
Index: []
```

```
In [60]: sample_submission['y'] = sample_submission['y'].map(lambda x: x if x > 0 else 0.0)
```

```
In [61]: sample_submission.to_csv("my_submission6.tsv", sep=',', index=False)
```